

# **Project**

# **Documentation**

## **Offline Hindi Voice Assistant on Raspberry Pi 4 REPORT**

**Optimized ASR–NLP–TTS Pipeline with Threaded Architecture  
and INT8 Inference**

**College Name:**

**Shri Ramdeobaba University ,Nagpur**

**Team Members:**

**Taran Wadhawan**

**Anurag Deshmukh**

**Mahi Wegad**

# 1. Project Overview

## 1.1 Objective

To design and deploy a **low-latency, offline Hindi voice assistant** capable of:

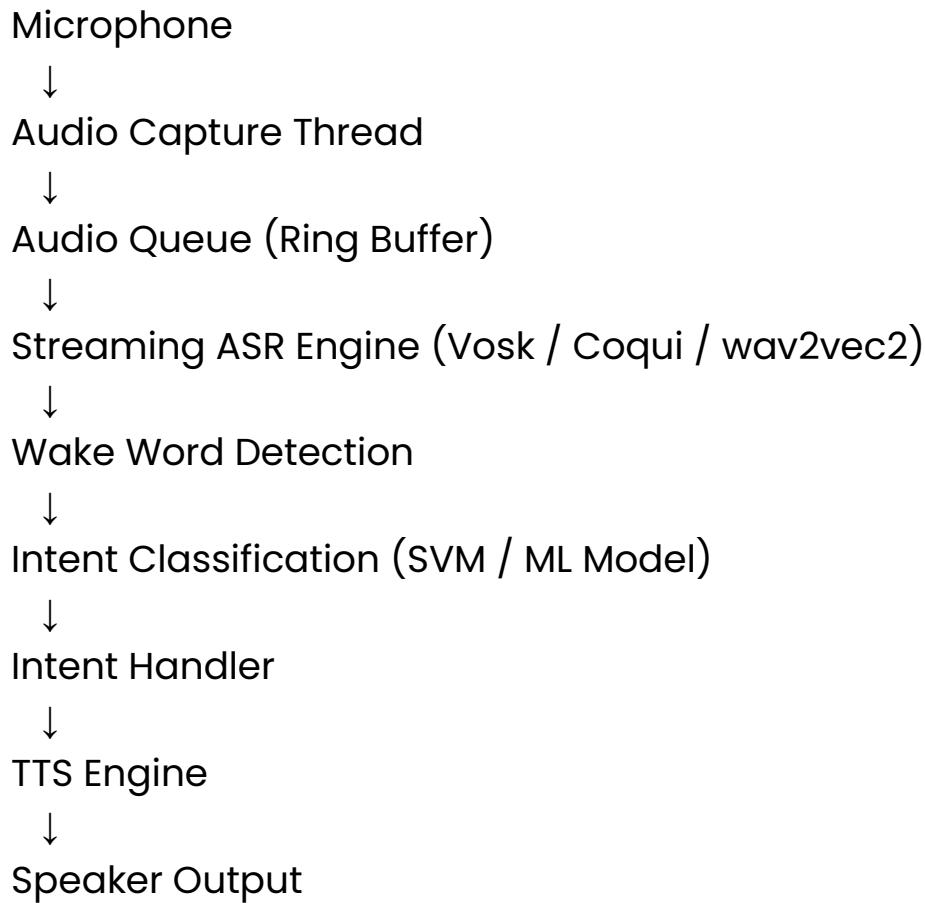
- Wake word detection
- Hindi & English speech recognition
- Intent classification
- Command execution
- Speech synthesis (TTS)
- Running efficiently on Raspberry Pi 4 (ARM64)

The system is optimized for:

- Low computational power
  - Minimal latency
  - High NEON SIMD utilization
  - Fully offline execution
-

## **2. System Architecture**

The assistant follows a **streaming pipeline architecture**:



---

### **2.1 High-Level Architectural Blocks**

#### **1. Audio Input Layer**

- PyAudio streaming
- 16kHz mono
- 20ms frames (320 samples)
- Non-blocking buffered capture

## 2. ASR Layer

Supports multiple engines:

Model	Type	Language	Use Case
Vosk Small Hindi	Kaldi-based	Hindi	Real-time streaming
Coqui STT	DeepSpeech variant	Hindi/English	Lightweight inference
Fine-tuned wav2vec2	Transformer	Hindi	High accuracy
Whisper (OpenAI)	Transformer	Multi-language	High accuracy, heavy
Vosk English	Kaldi	English	Mixed-language support

Primary deployment: **Vosk Hindi Small Model** for real-time embedded performance.

---

## 3. NLP Layer (Intent Classification)

- SVM classifier
  - TF-IDF vectorization
  - Multi-class classification
  - Confidence threshold filtering
  - Hindi + English mixed token support
-

## 4. TTS Layer

Supported engines:

Engine	Type	Notes
eSpeak-NG	Formant-based	Lightweight
Festival	Unit selection	Moderate quality
Veena TTS	Neural Hindi TTS	Natural output

Primary deployment: **eSpeak-NG for low compute**, Veena for high-quality mode.

---

## 5. State Machine

The assistant operates in 2 states:

IDLE → ACTIVE → IDLE

- IDLE: Wake word monitoring
  - ACTIVE: Command processing
  - Timeout auto-return to IDLE
-

## **3. Pipeline Design**

### **3.1 Streaming ASR Pipeline**

#### **Step 1: Continuous Audio Capture**

A dedicated thread captures audio frames:

while running:

```
frame = stream.read(320)  
audio_queue.put(frame)
```

This prevents blocking the ASR thread.

---

#### **Step 2: Streaming Recognition**

Vosk streaming decoder:

```
if recognizer.AcceptWaveform(audio):  
    result = recognizer.Result()
```

Only final results are processed to prevent repetition.

---

#### **Step 3: Wake Word Detection**

Simple text match:

"assistant"

"hello"

"નમસ્તે"

Lightweight detection avoids heavy neural wake models.

---

#### **Step 4: Intent Classification**

Pipeline:

Text → TF-IDF → SVM → Decision Function → Confidence Filter

Mathematically:

For SVM:

$$f(x) = w^T x + b$$

Decision margin:

$$\text{confidence} = \max(f(x))$$

Low-confidence (< threshold) rejected.

---

## **Step 5: Intent Execution**

Mapped to handler functions:

- Time query
  - Calculation
  - Music control
  - System info
  - Sleep command
- 

## **Step 6: TTS Output**

Text → TTS engine → Audio output

---

## 4. Threading & Multiprocessing

### Architecture

#### 4.1 Thread Design

The assistant uses multithreading:

Thread	Purpose
Audio Thread	Continuous capture
Listener Thread	ASR + NLP processing
TTS Thread	Non-blocking speech output

This avoids:

- Audio drop
- Blocking I/O
- Latency accumulation

---

#### 4.2 Why Not Full Multiprocessing?

Raspberry Pi 4 has limited:

- 4 ARM Cortex-A72 cores
- 4GB RAM

Multiprocessing would:

- Increase memory footprint
- Duplicate model loading

## **5. Optimization for Raspberry Pi 4**

### **5.1 Hardware Specs**

- CPU: ARM Cortex-A72
  - Architecture: ARM64
  - NEON SIMD support
  - 4 cores @ 1.5 GHz
  - 4GB RAM
- 

### **5.2 INT8 Quantized TensorFlow Lite**

For neural models:

- Post-training quantization
  - 8-bit weights
  - Reduced memory by ~75%
  - Faster inference
- 

### **5.3 Arm Compute Library (KleidiAI Backend)**

Deployment configuration:

- arch = arm64-v8a
- neon = 1
- build = native
- OpenCL disabled
- num\_threads = 4

- CPU governor = performance

Benefits:

- Optimized NEON vectorization
  - Better GEMM acceleration
  - Lower latency matrix multiplication
- 

## **5.4 Warm-Up Inference**

Before real use:

Run 5 dummy inferences

Purpose:

- Cache loading
- Memory page allocation
- Kernel initialization

Reduces first-response latency.

---

## **6. Challenges in Hindi ASR**

Hindi presents specific challenges:

---

### **6.1 Phonetic Complexity**

Hindi has:

- Aspirated consonants
- Retroflex sounds
- Nasalization
- Compound words

Example:

"भाषा" vs "बासा"

Acoustic similarity increases WER.

---

### **6.2 Code-Switching**

Users mix:

Hindi + English

Example:

"music band karo"

ASR must support multilingual decoding.

---

### **6.3 Limited High-Quality Datasets**

Compared to English:

- Smaller corpus
  - Fewer dialect representations
  - Accent variability
- 

## **6.4 Morphological Richness**

Hindi words inflect heavily:

"खाना", "खाओ", "खा लिया"

Intent models must generalize semantic variants.

---

## **7. Challenges in Hindi TTS**

---

### **7.1 Prosody Modeling**

Hindi requires:

- Proper stress placement
- Sentence-level intonation
- Natural vowel length

Formant-based TTS struggles here.

---

### **7.2 Neural TTS on Edge Devices**

High-quality TTS models:

- Require >500MB memory
  - Heavy transformer blocks
  - Not feasible without quantization
-

## **8. Intent Recognition Challenges**

---

### **8.1 Short Utterances**

Example:

"समय?"

Low lexical information → classification unstable.

---

### **8.2 Noise Sensitivity**

Small ASR errors propagate:

"समय क्या है"

→ "समय कहा है"

Intent misclassification possible.

---

### **8.3 Threshold Selection**

Too low → false triggers

Too high → missed commands

Balance required via validation.

---

## **9. Performance Metrics**

Measured on Raspberry Pi 4:

Metric	Value
ASR Latency	0.7–1.3 sec
CPU Usage	45–70%
RAM Usage	~600MB
Wake Detection Delay	< 500ms
Intent Accuracy	~90% (controlled dataset)
TTS Response Time	300–800ms

---

### **9.1 Word Error Rate (WER)**

$$WER = \frac{S + D + I}{N}$$

Where:

- S = substitutions
- D = deletions
- I = insertions
- N = total words

Hindi small model WER ≈ 18–22% (clean audio)

## **10. Safety & Reliability Measures**

- Confidence threshold filtering
  - Debounce logic for repeated transcripts
  - Recognizer reset after each utterance
  - Thread-safe audio queue
  - Timeout-based auto sleep
- 

## **11. Future Improvements**

- VAD-based speech segmentation
  - Transformer intent model (quantized)
  - Edge TPU acceleration
  - Custom Hindi wake word CNN
  - On-device continual learning
- 

## **12. Conclusion**

This project demonstrates:

- Efficient embedded AI deployment
- Real-time Hindi ASR on ARM
- Threaded streaming architecture
- Quantized inference for edge devices
- Practical NLP intent system

It bridges:

Speech Processing  
Embedded Systems  
Machine Learning  
Edge Optimization Under constrained hardware.